

Approximating Hessians for Neural Networks

Maciej Skorski

University of Luxembourg

May 24, 2021

Outline

- 1 Hessians in Neural Network Optimization
- 2 Approximate Chain Rule
- 3 Empirical Evaluation
- 4 Conclusion

Optimization Landscape of Neural Networks

- *curvature* is critical for training...
 - Ⓐ hessian controls accuracy of linearization (Taylor's formula)
 - Ⓑ preconditions step and convergence (even for the simplest SGD!)
 - Ⓒ used to boost convergence of Gradient Descents (AdaHess...)
 - Ⓓ used to initialize MC (NUTS)
- ... but hard to compute in high dimension, e.g. neural networks (*billions* and more of hessian entries)




Figure 1: LeNet5 loss (FashionMNIST). The dependency on > 44k parameters visualized along two random orthogonal directions.

Hessian Complexity

- Non-trivial back-propagation, different than for gradients ☹

$$D^2[f \circ g] = D^2f \bullet Dg \bullet Dg + Df \bullet D^2g,$$

where \circ and \bullet denote, resp., composition and tensordot.

- Memory-consuming and in general infeasible ☹
- Limited support in some of AD frameworks (JAX and TensorFlow support both forward and reverse mode but PyTorch not) ☹
- Limited knowledge resources - considered an advanced research topic.

Alternatives and Proposed Solution

- Compute HVPs (known in research, somewhat painful in practice)
 - Ⓐ theoretically fast ☺
 - Ⓑ limited in scope (vector products only) ☹
 - Ⓒ issues with AD-software support ☹
- Hutchinson's Trick (PyHess and others):
 - Ⓐ AD-software agnostic ☺
 - Ⓑ limited in scope (diagonal only) ☹
 - Ⓒ approximate and slow (probabilistic guarantees) ☹
- **Approximate Chain Rule** (this work)
 - Ⓐ agnostic of AD-software ☺
 - Ⓑ very fast ☺
 - Ⓒ unlimited in scope (all hessian parts) ☺
 - Ⓓ approximate, but works well for practical NNs! ☺

Approximate Chain Rule

- Let ℓ be the loss function and $z(w)$ be the network output as a function of the parameters w . Then the loss curvature is:

$$D_w^2 \ell(z(w)) = \underbrace{D_z^2 \ell(z) \bullet D_w z(w) \bullet D_w z(w)}_{\text{linearization effect}} + \underbrace{D_z \ell(z) \bullet D_w^2 z(w)}_{\text{curvature effect}}.$$

- For typical networks the curvature effect is of smaller order!
- We claim that with good accuracy

$$D_w^2 \ell(z(w)) \approx \underbrace{D_z^2 \ell(z) \bullet D_w z(w) \bullet D_w z(w)}_{\text{linearization effect}},$$

when *activations are nearly linear around 0* (satisfied in practice).

- The approximation costs just one back-propagation and two tensordots (hessian of loss wrt its natural domain is analytic) ☺

Local near-linearity of activations

- Approximation works for activations h s.t. $h(x) = \Theta(x) + O(x^3)$ for small x
- Satisfied for popular activations: linear, tanh, sigmoid, ReLU




Figure 2: Behavior of activations around 0: second Taylor's term vanishes!

- This allow us to prove that the "curvature contribution" is negligible

Chain Rule Evaluation

- error term is of smaller order (theoretical guarantees)
- *no error* with variations of ReLU
- *small error* on popular nets: LeNet, EfficientNet, ResNet...




Figure 3: Approximation error at init. Main hessian components are evaluated on gradients.

Applications

- Quick approximate evaluation of any hessian components
- Speeding-up hessian-based optimizers
- Proofs for init schemes (replacing "variance-flow" heuristics)
- Improving convergence of specific nets, see the case study below ☺

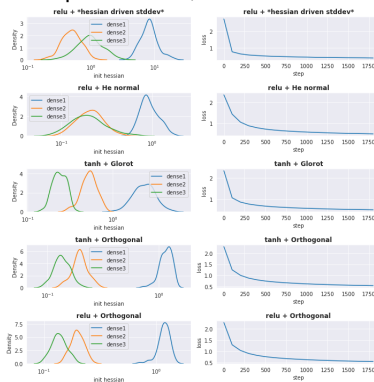


Figure 4: Network with layers of 128, 84 and 10 units on FashionMNIST. Note that the hessian-driven initialization outperforms established initializers.

Summing Up...

- We have a fast and accurate chain rule for hessian of NNs
- The idea is to neglect "curvature effects" appearing due to nearly-linear activation functions
- The approximation tested on a range of architectures
- Code at https://github.com/maciejskorski/ml_examples!



Working Philosophy ☺

- *learn diversity* of techniques (calculus, linear algebra, statistics)
- *divide and conquer*: attack problems with established tools
- *bridge to applications*: abstract findings, link to other areas...




Figure 5: Love this diagram!